

Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Step 1: Log in to your AWS Academy/personal account and launch 3 new Ec2 Instances. Select Ubuntu as AMI and t2.medium as Instance Type and create a key of type RSA with .pem extension and move the downloaded key to the new folder. We can use 3 Different keys or 1 common key also. Note: A minimum of 2 CPUs are required so Please select t2.medium and do not forget to stop the instance after the experiment because it is not available in the free tier.

Master:

Name and tags [Info](#)

Name: master [Add additional tags](#)

▼ Application and OS Images (Amazon Machine Image) [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Search our full catalog including 1000s of application and OS images

Recents **Quick Start**

Amazon Linux, macOS, Ubuntu, Windows, Red Hat, SUSE Linux

▼ Instance type [Info](#) [Get advice](#)

Instance type: t2.medium

Family: t2 - 2 vCPU - 4 GiB Memory - Current generation: true

On-Demand Linux base pricing: 0.0464 USD per Hour

On-Demand RHEL base pricing: 0.0752 USD per Hour

On-Demand Windows base pricing: 0.0644 USD per Hour

On-Demand SUSE base pricing: 0.1464 USD per Hour

☒ All generations [Compare instance types](#)

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required: Rakshit [Create new key pair](#)

Security group rule 4 (TCP, b443) [Remove](#)

▼ Security group rule 5 (TCP, 10251) [Remove](#)

Type: Custom TCP [Info](#)

Protocol: TCP [Info](#)

Port range: 10251 [Info](#)

Source type: Custom [Info](#)

Source: [Add CIDR, prefix list or security group](#)

Description - optional: e.g. SSH for admin desktop [Info](#)

Security group rule 6 (TCP, 10250) [Remove](#)

Type: Custom TCP [Info](#)

Protocol: TCP [Info](#)

Port range: 10250 [Info](#)

Source type: Custom [Info](#)

Source: [Add CIDR, prefix list or security group](#)

Description - optional: e.g. SSH for admin desktop [Info](#)

Security group rule 7 (TCP, 10252) [Remove](#)

Type: Custom TCP [Info](#)

Protocol: TCP [Info](#)

Port range: 10252 [Info](#)

Source type: Custom [Info](#)

Source: [Add CIDR, prefix list or security group](#)

Description - optional: e.g. SSH for admin desktop [Info](#)

Worker:

The screenshot shows the AWS Management Console interface. On the left, the 'Application and OS Images (Amazon Machine Image)' section is visible, showing a search bar and a grid of AMIs including Amazon Linux, macOS, Ubuntu, Windows, Red Hat, and SUSE Linux. The 'Quick Start' tab is selected. On the right, the 'Security groups' section is visible, showing three security groups: 'Security group rule 4 (TCP: 10250, 0.0.0.0/0)', 'Security group rule 5 (TCP: 30000-32767, 0.0.0.0/0)', and 'Security group rule 6 (TCP: 0-65535)'. Each rule has a 'Type' of 'Custom TCP', a 'Protocol' of 'TCP', and a 'Port range' of '10250', '30000-32767', and '0-65535' respectively. The 'Source type' is 'Custom' and the 'Source' is '0.0.0.0/0'.

Step 2: After creating the instances click on Connect & connect all 3 instances and navigate to SSH Client.

The screenshot shows the AWS Management Console interface. The top section is the 'Instances' page, showing a list of instances with columns for Name, Instance ID, Instance state, Instance type, Status check, and Alarm status. The instances are: 'nagios-host' (Terminated), 'master' (Running), 'node' (Running), and 'node' (Running). The 'master' instance is selected. Below the list, the 'Connect to instance' page is shown, with tabs for 'EC2 Instance Connect', 'Session Manager', 'SSH client', and 'EC2 serial console'. The 'SSH client' tab is selected. The page shows the instance ID 'i-0a9a7aa3d967ebcd8 (node2)' and provides instructions for connecting to the instance using an SSH client. The instructions include: 1. Open an SSH client. 2. Locate your private key file. The key used to launch this instance is Rakshit.pem. 3. Run this command, if necessary, to ensure your key is not publicly viewable. 4. Connect to your instance using its Public DNS: ec2-54-197-87-163.compute-1.amazonaws.com. An example command is provided: `ssh -i "Rakshit.pem" ubuntu@ec2-54-197-87-163.compute-1.amazonaws.com`. A note states: 'Note: In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.'

Step 3: Now open the folder in the terminal 3 times for Master, Node1& Node 2 where our .pem key is stored and paste the Example command (starting with ssh -i) in the terminal.(ssh -i "Master_Ec2_Key.pem" ubuntu@ec2-54-196-129-215.compute-1.amazonaws.com) Master:

```
Expanded Security Maintenance for Applications is not enabled.
```

```
0 updates can be applied immediately.
```

```
Enable ESM Apps to receive additional future security updates.  
See https://ubuntu.com/esm or run: sudo pro status
```

```
The programs included with the Ubuntu system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/*/copyright.
```

```
Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by  
applicable law.
```

```
To run a command as administrator (user "root"), use "sudo <command>".  
See "man sudo_root" for details.
```

```
ubuntu@ip-172-31-81-239:~$ |
```

Step 4: Run on Master, Node 1, and Node 2 the below commands to install and setup Docker in Master, Node1, and Node2.

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add - curl -fsSL  
https://download.docker.com/linux/ubuntu/gpg | sudo tee /etc/apt/trusted.gpg.d/docker.gpg >  
/dev/null sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu  
$(lsb_release -cs) stable"
```

```
ubuntu@ip-172-31-81-239:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -  
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee  
/etc/apt/trusted.gpg.d/docker.gpg > /dev/null  
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu  
$(lsb_release -cs) stable"  
Warning: apt-key is deprecated. Manage keyring files in trusted.gpg.d instead (see apt-key(8)).  
OK  
-----BEGIN PGP PUBLIC KEY BLOCK-----  
  
mQINBFit2ioBEADhWpZ8/wvZ6hUTiX0wQHXMALaFhcPH9hAtr4F1y2+0YdbtMuth  
lqqwp028AqyY+PRFVMtSYMBjuQuu5byyKR01BbqYhuS3jtqQmLjZ/bJvXqnmVXh  
38UuLa+z077PxyxQhu5BbqntTPQMfiyqEiU+BKbq2WmANUKQf+1AmZY/IruOXbnq  
L4C1+gJ8vfmXQt99npCaxEjaNRVYfOS8QcixNzHUYnbn6emjLANyEVLZzeqo7XKL7  
UrwV5inawTSzWnvtjEjj4nJL8NsLwscpLPQUhTQ+7BbQXAwAmeHCUTQIvvWXqw0N  
cmhh4HgeQscQHYgOJjjDVfoY5MucvglbIgCqfzAHW9jxmRL4qbMZj+b1XoePEtht  
ku4bIQN1X5P07fNwzlgaRL5Z4POXDDZTLIQ/EL58j9kp4bnWRCJW0lya+f8ocodo  
vZZ+Doi+fy4D5ZGrL4XEcIQP/Lv5uFyf+kQtL/94VFYVJ0leAv8W92KdgDkhTcTD  
G7c0tIkVEKNUq48b3aQ64NOZQW7fVjfoKwEZdOqPE72Pa45jrZzvUFxSpdiNk2tZ  
XYukHjlxxEgBdC/J3cMMNRE1F4NCA3ApfV1Y7/hTeOnmDuDYwr9/obA8t016YLjj  
q5rdkywPf4JF8mXUW5eCN1vAFHxeg9ZWemhBtQmGxXnw9M+z6hWwc6ahmwARAQAB  
+CtEb2NrZXIuLmVzZS4oO0UuZGVvYyYyKR01BbqYhuS3jtqQmLjZ/bJvXqnmVXh
```

```
sudo apt-get update
```

sudo apt-get install -y docker-ce

```
ubuntu@ip-172-31-81-239:~$ sudo apt-get update
sudo apt-get install -y docker-ce
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Hit:5 http://security.ubuntu.com/ubuntu noble-security InRelease
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  containerd.io docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras docker-compose-plugin libltdl7 libslirp0
  pigz slirp4netns
Suggested packages:
```

sudo mkdir -p /etc/docker

cat <<EOF | sudo tee /etc/docker/daemon.json

```
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
```

```
ubuntu@ip-172-31-81-239:~$ sudo mkdir -p /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
```

sudo systemctl enable docker

sudo systemctl daemon-reload

sudo systemctl restart docker

```
ubuntu@ip-172-31-81-239:~$ sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
Synchronizing state of docker.service with SysV service script with /usr/lib/systemd/systemd-sysv-install.
Executing: /usr/lib/systemd/systemd-sysv-install enable docker
```

Step 5: Run the below command to install Kubernetes.

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg echo 'deb
[signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
```

`https://pkgs.k8s.io/core:/stable:/v1.31/deb/ '/' | sudo tee /etc/apt/sources.list.d/kubernetes.list`

```
ubuntu@ip-172-31-81-239:~$ curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ '/' | sudo tee /etc/apt/sources.list.d/kubernetes.list
gpg: missing argument for option "-o"
-bash: /etc/apt/keyrings/kubernetes-apt-keyring.gpg: No such file or directory
deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /
```

`sudo apt-get update`

`sudo apt-get install -y kubelet kubeadm kubectl`

`sudo apt-mark hold kubelet kubeadm kubectl`

`sudo systemctl enable --now kubelet`

`sudo apt-get install -y containerd`

```
ubuntu@ip-172-31-81-239:~$ sudo apt-get install -y apt-transport-https ca-certificates curl
curl -fsSL https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -
echo "deb https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee /etc/apt/sources.list.d/kubernetes.list
sudo apt-get update
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
ca-certificates is already the newest version (20240203).
ca-certificates set to manually installed.
curl is already the newest version (8.5.0-2ubuntu10.4).
curl set to manually installed.
The following NEW packages will be installed:
  apt-transport-https
0 upgraded, 1 newly installed, 0 to remove and 6 not upgraded.
Need to get 3974 B of archives.
After this operation, 35.8 kB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 apt-transport-https all 2.7.14build2 [3974 B]
Fetched 3974 B in 0s (605 kB/s)
debconf: delaying package configuration, since apt-utils is not installed
```

`sudo mkdir -p /etc/containerd`

`sudo containerd config default | sudo tee /etc/containerd/config.toml`

```
ubuntu@ip-172-31-81-239:~$ sudo mkdir -p /etc/containerd
sudo containerd config default | sudo tee /etc/containerd/config.toml
disabled_plugins = []
imports = []
oom_score = 0
plugin_dir = ""
required_plugins = []
root = "/var/lib/containerd"
state = "/run/containerd"
temp = ""
version = 2

[cgroup]
  path = ""

[debug]
  address = ""
  format = ""
  gid = 0
  level = ""
```

`sudo systemctl restart containerd`

`sudo systemctl enable containerd`

sudo systemctl status containerd

```
ubuntu@ip-172-31-81-239:~$ sudo systemctl restart containerd
sudo systemctl enable containerd
sudo systemctl status containerd
● containerd.service - containerd container runtime
   Loaded: loaded (/usr/lib/systemd/system/containerd.service; enabled; preset: enabled)
   Active: active (running) since Thu 2024-10-03 12:03:56 UTC; 187ms ago
     Docs: https://containerd.io
   Main PID: 4388 (containerd)
      Tasks: 8
     Memory: 13.8M (peak: 14.1M)
        CPU: 47ms
    CGroup: /system.slice/containerd.service
            └─4388 /usr/bin/containerd

Oct 03 12:03:56 ip-172-31-81-239 containerd[4388]: time="2024-10-03T12:03:56.810259144Z" level=info msg="Start subscrib
Oct 03 12:03:56 ip-172-31-81-239 containerd[4388]: time="2024-10-03T12:03:56.810436084Z" level=info msg="Start recoveri
Oct 03 12:03:56 ip-172-31-81-239 containerd[4388]: time="2024-10-03T12:03:56.810483329Z" level=info msg="servin... addr
Oct 03 12:03:56 ip-172-31-81-239 containerd[4388]: time="2024-10-03T12:03:56.810489545Z" level=info msg="Start event mo
Oct 03 12:03:56 ip-172-31-81-239 containerd[4388]: time="2024-10-03T12:03:56.810583963Z" level=info msg="Start snapshot
Oct 03 12:03:56 ip-172-31-81-239 containerd[4388]: time="2024-10-03T12:03:56.810599063Z" level=info msg="Start cni netw
Oct 03 12:03:56 ip-172-31-81-239 containerd[4388]: time="2024-10-03T12:03:56.810605397Z" level=info msg="Start streamin
Oct 03 12:03:56 ip-172-31-81-239 containerd[4388]: time="2024-10-03T12:03:56.810635592Z" level=info msg="servin... addr
Oct 03 12:03:56 ip-172-31-81-239 systemd[1]: Started containerd.service - containerd container runtime.
Oct 03 12:03:56 ip-172-31-81-239 containerd[4388]: time="2024-10-03T12:03:56.813064087Z" level=info msg="containerd suc
```

sudo apt-get install -y socat

```
ubuntu@ip-172-31-81-239:~$ sudo apt-get install -y socat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras docker-compose-plugin libltdl7 libslirp0 pigz
  slirp4netns
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  socat
0 upgraded, 1 newly installed, 0 to remove and 6 not upgraded.
Need to get 374 kB of archives.
After this operation, 1649 kB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 socat 1.8.0.0-4build3 [374 kB]
Fetched 374 kB in 0s (12.2 MB/s)
Selecting previously unselected package socat.
(Reading database ... 68148 files and directories currently installed.)
Preparing to unpack .../socat_1.8.0.0-4build3_amd64.deb ...
Unpacking socat (1.8.0.0-4build3) ...
Setting up socat (1.8.0.0-4build3) ...
```

Step 6: Initialize the Kubecuster .Now Perform this Command only for Master.

sudo kubeadm init --pod-network-cidr=10.244.0.0/16

```
ubuntu@ip-172-31-82-119:~$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
W1003 12:37:31.290835 10158 checks.go:846] detected that the sandbox image "registry.k8s.io/pause:3.8" of the containe
r runtime is inconsistent with that used by kubeadm.It is recommended to use "registry.k8s.io/pause:3.10" as the CRI san
dbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-82-119 kubernetess kubernetess.default kubernetess.default
t.svc kubernetess.default.svc.cluster.local] and IPs [10.96.0.1 172.31.82.119]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-172-31-82-119 localhost] and IPs [172.31.82.119 127.0.0.1 :
:1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-172-31-82-119 localhost] and IPs [172.31.82.119 127.0.0.1 ::1
```

Run this command on master and also copy and save the Join command from above.

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.82.119:6443 --token h4iisg.g9tdfmc88m9toefp \
--discovery-token-ca-cert-hash sha256:5b3c7cfd8115a26f1e73b752820d2b27b84ed476b344aaa3bd1a90e4b1f2105
ubuntu@ip-172-31-82-119:~$ mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Step 7: Now Run the command `kubectl get nodes` to see the nodes before executing Join command on nodes.

```
ubuntu@ip-172-31-82-119:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
ip-172-31-82-119    NotReady control-plane 82s   v1.31.1
```

Step 8: Now Run the following command on Node 1 and Node 2 to Join to master.

```
sudo kubeadm join 172.31.82.119:6443 --token h4iisg.g9tdfmc88m9toefp \
--discovery-token-ca-cert-hash
sha256:5b3c7cfd8115a26f1e73b752820d2b27b84ed476b344aaa3bd1a90e4b1f2105
```

```
ubuntu@ip-172-31-84-169:~$ sudo kubeadm join 172.31.82.119:6443 --token h4iisg.g9tdfmc88m9toefp \
--discovery-token-ca-cert-hash sha256:5b3c7cfd8115a26f1e73b752820d2b27b84ed476b344aaa3bd1a90e4b1f2105
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 501.315929ms
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

Step 9: Now Run the command `kubectl get nodes` to see the nodes after executing Join command on nodes.

```
ubuntu@ip-172-31-82-119:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ip-172-31-82-119	NotReady	control-plane	2m56s	v1.31.1
ip-172-31-84-169	NotReady	<none>	43s	v1.31.1
ip-172-31-87-189	NotReady	<none>	39s	v1.31.1

Step 10: Since Status is NotReady we have to add a network plugin. And also we have to give the name to the nodes.

kubectl apply -f <https://docs.projectcalico.org/manifests/calico.yaml>

```
ubuntu@ip-172-31-82-119:~$ kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
poddisruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
```

sudo systemctl status kubelet

```
ubuntu@ip-172-31-82-119:~$ sudo systemctl status kubelet
● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/usr/lib/systemd/system/kubelet.service; enabled; preset: enabled)
   Drop-In: /usr/lib/systemd/system/kubelet.service.d
            └─10-kubeadm.conf
   Active: active (running) since Thu 2024-10-03 12:37:52 UTC; 4min 4s ago
     Docs: https://kubernetes.io/docs/
   Main PID: 10849 (kubelet)
    Tasks: 10 (limit: 4676)
   Memory: 32.3M (peak: 33.0M)
      CPU: 5.030s
   CGroup: /system.slice/kubelet.service
           └─10849 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/k

Oct 03 12:41:52 ip-172-31-82-119 kubelet[10849]: E1003 12:41:52.678397 10849 kuberuntime_container.go:851] "Kill cont
Oct 03 12:41:52 ip-172-31-82-119 kubelet[10849]:         rpc error: code = Unknown desc = failed to kill container "233
Oct 03 12:41:52 ip-172-31-82-119 kubelet[10849]:         : unknown
Oct 03 12:41:52 ip-172-31-82-119 kubelet[10849]: > pod="kube-system/etcd-ip-172-31-82-119" podUID="201366195a044a56173
Oct 03 12:41:52 ip-172-31-82-119 kubelet[10849]: E1003 12:41:52.688038 10849 log.go:32] "StopPodSandbox from runtime >
Oct 03 12:41:52 ip-172-31-82-119 kubelet[10849]:         rpc error: code = Unknown desc = failed to stop container "233
Oct 03 12:41:52 ip-172-31-82-119 kubelet[10849]:         : unknown
```


Now Run command `kubectl get nodes -o wide` we can see Status is ready.

```
ubuntu@ip-172-31-82-119:~$ kubectl get nodes -o wide
NAME                                STATUS    ROLES    AGE     VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-
VERSION   CONTAINER-RUNTIME
ip-172-31-82-119 Ready     control-plane 4m32s   v1.31.1   172.31.82.119 <none>        Ubuntu 24.04.1 LTS   6.8.0-1
016-aws   containerd://1.7.12
ip-172-31-84-169 Ready     <none>      2m19s   v1.31.1   172.31.84.169 <none>        Ubuntu 24.04.1 LTS   6.8.0-1
016-aws   containerd://1.7.12
ip-172-31-87-189 Ready     <none>      2m15s   v1.31.1   172.31.87.189 <none>        Ubuntu 24.04.1 LTS   6.8.0-1
016-aws   containerd://1.7.12
```

Now to Rename run this command `kubectl label node ip-172-31-18-135`

`kubernetes.io/role=worker` Rename to Node 1:`kubectl label node ip-172-31-28-117`

`kubernetes.io/role=Node1` Rename to Node 2:`kubectl label node ip-172-31-18-135`

`kubernetes.io/role=Node2`

Step 11: Run command `kubectl get nodes -o wide` . And Hence we can see we have Successfully connected Node 1 and Node 2 to the Master.

```
ubuntu@ip-172-31-82-119:~$ kubectl get nodes
NAME                                STATUS    ROLES    AGE     VERSION
ip-172-31-82-119 Ready     control-plane 6m52s   v1.31.1
ip-172-31-84-169 Ready     Node1      4m39s   v1.31.1
ip-172-31-87-189 Ready     Node2      4m35s   v1.31.1
```

Conclusion:

In this experiment, we successfully set up a Kubernetes cluster on AWS EC2 instances. We began by provisioning three EC2 instances—Master, Node1, and Node2—with appropriate security rules to facilitate communication within the cluster. After securely SSHing into each instance, we installed Docker as the container runtime on all machines. We proceeded by installing Kubernetes on each instance and initializing the cluster from the master node. Following cluster initialization, the worker nodes were connected to the master using the `kubeadm join` command, effectively forming a cohesive Kubernetes cluster. Additionally, we implemented a network plugin to ensure proper communication between nodes, allowing the nodes' status to transition to "Ready."