
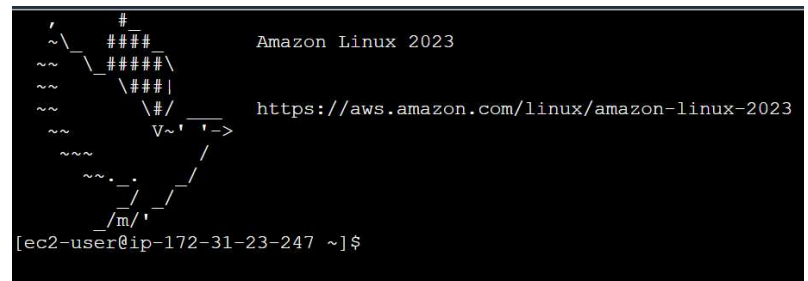
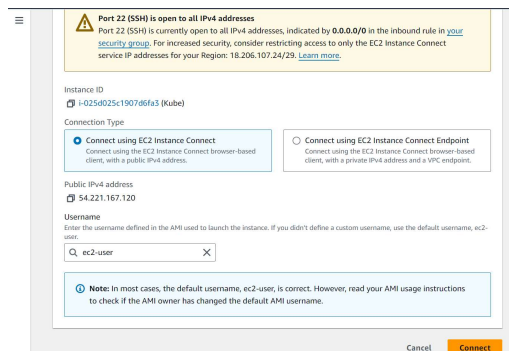


**Aim:** To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy Your First Kubernetes Application.

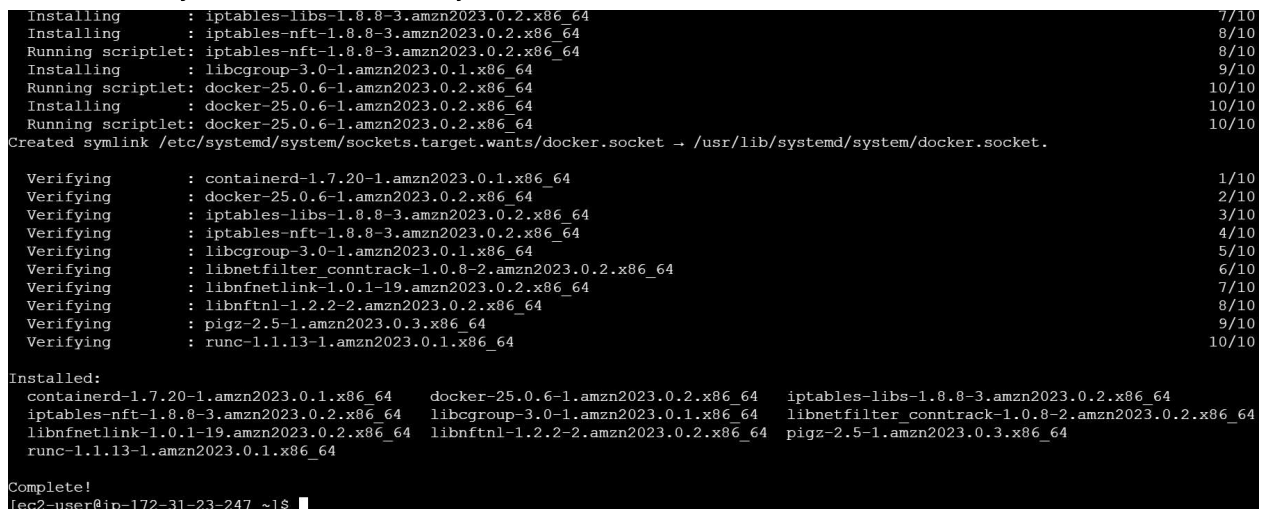
1. Create an EC2 instance with OS as Amazon Linux and make sure to allow SSH traffic.



2. Select the instance and click on connect. On the page scroll down and click on connect to open the command line.



3. To install docker run the following command:  
`sudo yum install docker -y`



4. Configure cgroup in daemon.json file using the following commands:

```
cd /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
```

```
[ec2-user@ip-172-31-23-247 ~]$ cd /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
```

5. Run the following command after this:

```
sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
```

```
[ec2-user@ip-172-31-23-247 docker]$ sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[ec2-user@ip-172-31-23-247 docker]$
```

6. Verify installation using docker -v command

```
[ec2-user@ip-172-31-23-247 docker]$ docker -v
Docker version 25.0.5, build 5dc9bcc
```

7. Install Kubernetes

I. Disable SELinux before configuring kubelet

```
sudo setenforce 0
```

```
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/'  
/etc/selinux/config
```

```
[ec2-user@ip-172-31-23-247 docker]$ sudo setenforce 0  
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config  
[ec2-user@ip-172-31-23-247 docker]$
```

II. Add kubernetes repository

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
```

```
[kubernetes]
```

```
name=Kubernetes
```

```
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
```

```
enabled=1
```

```
gpgcheck=1
```

```
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repo  
data/repomd.xml.key
```

```
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
```

```
EOF
```

```
[ec2-user@ip-172-31-23-247 docker]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo  
[kubernetes]  
name=Kubernetes  
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/  
enabled=1  
gpgcheck=1  
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repo  
data/repomd.xml.key  
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni  
EOF  
[kubernetes]  
name=Kubernetes  
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/  
enabled=1  
gpgcheck=1  
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repo  
data/repomd.xml.key  
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni  
[ec2-user@ip-172-31-23-247 docker]$
```

III. Run the commands to update and install kubernetes packages

```
sudo yum update
```

```
sudo yum install -y kubelet kubeadm kubectl
```

```
--disableexcludes=kubernetes
```

```

Installing      : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64      4/9
Installing      : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64      5/9
Installing      : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64             6/9
Running scriptlet: conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64             6/9
Installing      : kubelet-1.31.1-150500.1.1.x86_64                       7/9
Running scriptlet: kubelet-1.31.1-150500.1.1.x86_64                       7/9
Installing      : kubeadm-1.31.1-150500.1.1.x86_64                      8/9
Installing      : kubectrl-1.31.1-150500.1.1.x86_64                     9/9
Running scriptlet: kubectrl-1.31.1-150500.1.1.x86_64                     9/9
Verifying       : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64             1/9
Verifying       : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64      2/9
Verifying       : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64     3/9
Verifying       : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64         4/9
Verifying       : cri-tools-1.31.1-150500.1.1.x86_64                    5/9
Verifying       : kubeadm-1.31.1-150500.1.1.x86_64                     6/9
Verifying       : kubectrl-1.31.1-150500.1.1.x86_64                     7/9
Verifying       : kubelet-1.31.1-150500.1.1.x86_64                     8/9
Verifying       : kubernetes-cni-1.5.1-150500.1.1.x86_64                9/9

Installed:
  conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64      cri-tools-1.31.1-150500.1.1.x86_64
  kubeadm-1.31.1-150500.1.1.x86_64                kubectrl-1.31.1-150500.1.1.x86_64
  kubelet-1.31.1-150500.1.1.x86_64                kubernetes-cni-1.5.1-150500.1.1.x86_64
  libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64
  libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64   libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64

Complete!
[ec2-user@ip-172-31-23-247 docker]$

```

#### IV. Configure internet options to allow bridging

- `sudo swapoff -a`
- `echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf`
- `sudo sysctl -p`

```

[ec2-user@ip-172-31-23-247 docker]$ sudo swapoff -a
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
sudo sysctl -p
net.bridge.bridge-nf-call-iptables=1
net.bridge.bridge-nf-call-iptables = 1

```

#### 8. Initialize the kubecoluster

`sudo kubeadm init --pod-network-cidr=10.244.0.0/16`

```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectrl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.23.247:6443 --token x667nq.2cid7otgtlazgufa \
  --discovery-token-ca-cert-hash sha256:84137c58548c42038b77cc5f5a59847bd054230f5f80b0153067884c534d980c

```

Save the join command in notepad as it will be used later.

Run the 3 commands starting from mkdir given above.

```
[ec2-user@ip-172-31-23-247 docker]$ mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
[ec2-user@ip-172-31-23-247 docker]$
```

Add a common network plugin called Flannel as mentioned in the code below:

kubectl apply -f

<https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>

```
c2-user@ip-172-31-20-245 dockekubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
mespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
c2-user@ip-172-31-20-245 docker]$
```

Cluster is up and running

9. Deploy nginx server on this cluster using the command

kubectl apply -f <https://k8s.io/examples/pods/simple-pod.yaml>

Also run kubectl get pods to check creation of pod

```
[ec2-user@ip-172-31-20-245 ~]$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
nginx         0/1     Pending   0           80s
```

To change the state from pending to running, use the following command

kubectl describe pod nginx

This command will help to describe the pods it gives reason for failure as it shows the untolerated taints which need to be untainted.

```
Containers:
  nginx:
    Image:          nginx:1.14.2
    Port:           80/TCP
    Host Port:      0/TCP
    Environment:    <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-dmncs (ro)
Conditions:
  Type            Status
  PodScheduled    False
Volumes:
  kube-api-access-dmncs:
    Type:          Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:    kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI:      true
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                 node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type            Reason              Age             From              Message
  ----            -
  Warning         FailedScheduling    2m47s          default-scheduler  0/1 nodes are available: 1 node(s) had untolerated taint (node-role.kubernetes.io/control-plane: ). Preemption: 0/1 nodes are available: 1 Preemption is not helpful for scheduling.
```

```
[ec2-user@ip-172-31-20-245 docker]$ kubectl taint nodes ip-172-31-20-245.ec2.internal node-role.kubernetes.io/control-plane-
node/ip-172-31-20-245.ec2.internal untainted
```

## 10. Check the status of pod

```
[ec2-user@ip-172-31-20-245 docker]$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	4m3s

## 11. Lastly, mention the port you want to host. Here I have used localhost 8081 then check it.

```
kubectl port-forward nginx 8081:80
```

```
[ec2-user@ip-172-31-20-245 docker]$ kubectl port-forward nginx 8081:80
Forwarding from 127.0.0.1:8081 -> 80
Forwarding from [::1]:8081 -> 80
```

**Conclusion:**

In this experiment, we successfully configured a Kubernetes environment on an Amazon Linux EC2 instance. We installed Docker and adjusted its settings to use `systemd` for cgroup management. We then set up Kubernetes by disabling SELinux, configuring the Kubernetes repository, and installing necessary components. After initializing the cluster and applying the Flannel network plugin, we deployed an Nginx server. We also addressed issues related to pod scheduling and port forwarding, ensuring the Nginx pod was accessible via port 8081 on local machine.